

# BITS F464 Machine Learning

## Multi-Layer Perceptron

### 1. Introduction:

In this report, we try to analyse the performance (accuracy, fscore, recall and precision) of a neural network that we built using NumPy. The implementation of this network is similar to that of Keras and the whole code has been written in a completely modularized fashion.

### 2. Data Preprocessing:

The dataset *housepricesdata.csv* contained 1460 data points, each containing 10 input features (representing house features) and one binary target variable (representing if the house price is above or below the median price). In the preprocessing stage, all features were standardized and randomly split into *train\_set* and *test\_set* (80-20 split ratio).

### 3. Activation functions:

- linear:  $y = x$
- tanh:  $y = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
- sigmoid:  $y = \frac{1}{1 + e^{-x}}$
- ReLU:  $y = \max(0, x)$

### 4. Loss function:

- $L_n = -t_n \log(y_n) - (1 - t_n) \log(1 - y_n)$

### 5. Network Architecture:

#### 5.1 Two-Layer NN

The first model we built was a 2-layer neural network (i.e. with one hidden layer). We experimented with activations like linear, ReLU and tanh for the hidden layer. The input layer had 10 neuron units with a ReLU activation (for best results. A slight drop in accuracy was observed when linear activation was used for the input layer). The final layer had an activation of sigmoid for binary classification. Output value greater than or equal to 0.5 was taken to represent positive classes and less than 0.5 to represent negative classes. `np.random.random()` and `np.random.randn()` methods were used to randomly initialize weights and biases associated network to initialize the weights according to a Uniform or Guassian distribution respectively. Best results were obtained using uniformly distributed weights/biases with a seed=5 and 4 hidden neurons with ReLU and learning\_rate=0.001 with 20000 training epochs. The loss was observed to have more fluctuations using this configuration after 40000 epochs and thus lowering the accuracy. The avg. loss using this configuration was observed to be 0.7 (per data point). The highest accuracy achieved by this model using this configuration was 87.67% and the lowest accuracy was 80% with an average accuracy of 84%.

## 5.2 Three-Layer NN

The second model we built and experimented with was a three layer neural network. The accuracy of the three layer neural network was found to be better than that of the two layer network. Both the hidden layers had 5 neurons with ReLU activation for optimal results. Training the network, in this case, required more iterations as there were more number of parameter to be tuned than the two layer network. The optimal number of neurons for both hidden layers were chosen accordingly as suggested in research article *Review on Methods to Fix Number of Hidden Neurons in Neural Networks* by Gnana Sheela et al. Once again, best results were obtained using uniform distribution of weights with learning\_rate=0.01 and 40000 iterations. The following combination yielded the highest accuracy of 91.43% with a minimum accuracy of 80% and an average accuracy of 86%.

## 6. Conclusion:

- In general, a three layer network performs better than a two layer network. Adding more layers/more neurons to the hidden layers provides more trainable parameters which help the network to fit the model in a better way. But the drawback of adding multiple layers/multiple neurons is that, the model requires more iterations to train. In this case, there is a higher chance that the model might overfit the data, and hence giving a low testing accuracy.
- A lower learning rate requires more iterations to converge, but conversely, using a higher learning rate might lead to fluctuation in the loss/exploding gradient.
- Different initialization of weights leads to convergence of the model at different local minima (initializations mentioned above).
- Different activation functions for the hidden layers provides drastically varying results. We have used ReLU (Rectified Linear) which is by and the far, the de-facto standard activation used for hidden layers. Other activations like tanh or sigmoid gave very low accuracy as compared to that of ReLU. ReLU suffers from the *dying neuron problem*, in which the output of a neuron might be 0 for most of the training (i.e. that neuron is not contributing to the network), and hence we tried using leaky-ReLU, but the accuracy was found to more or less the same.

[results on next page]

## 7. Additional Work:

- **SGD Momentum and Decay**

$$M_t = \beta M_{t-1} + (1 - \beta) \nabla E$$
$$W = W - \eta M_t$$

- **RMSprop**

$$V_t = \beta V_{t-1} + (1 - \beta) (\nabla E)^2$$
$$W = W - \eta \frac{\nabla E}{\sqrt{V_t} + \epsilon}$$

- **Adam**

$$M_t = \beta M_{t-1} + (1 - \beta_1) \nabla E$$
$$V_t = \beta V_{t-1} + (1 - \beta_2) (\nabla E)^2$$
$$\hat{M}_t = \frac{M_t}{1 - \beta_1^t}; \hat{V}_t = \frac{V_t}{1 - \beta_2^t}$$
$$W = W - \eta \frac{\hat{M}_t}{\sqrt{\hat{V}_t} + \epsilon}$$

## 8. Results:

<pre>training metrics: tp = 508, tn = 479, fp = 111, fn = 70 final accuracy: 84.50342465753424 final recall: 87.88927335640139 final precision: 82.06785137318255 final fscore: 84.87886382623225  testing metrics: tp = 133, tn = 123, fp = 19, fn = 17 final accuracy: 87.67123287671232 final recall: 88.66666666666667 final precision: 87.5 final fscore: 88.0794701986755</pre>	<pre>training metrics: tp = 544, tn = 402, fp = 188, fn = 34 final accuracy: 80.9931506849315 final recall: 94.11764705882352 final precision: 74.31693989071039 final fscore: 83.05343511450381  testing metrics: tp = 144, tn = 107, fp = 35, fn = 6 final accuracy: 85.95890410958904 final recall: 96.0 final precision: 80.44692737430168 final fscore: 87.53799392097265</pre>
---	--

### 2-Layer NN using uniform distribution

<pre>training metrics: tp = 552, tn = 453, fp = 137, fn = 26 final accuracy: 86.0445205479452 final recall: 95.50173010380622 final precision: 80.11611030478954 final fscore: 87.13496448303077  testing metrics: tp = 143, tn = 111, fp = 31, fn = 7 final accuracy: 86.98630136986301 final recall: 95.33333333333334 final precision: 82.18390804597702 final fscore: 88.27160493827162</pre>	<pre>training metrics: tp = 521, tn = 495, fp = 95, fn = 57 final accuracy: 86.98630136986301 final recall: 90.138408304449827 final precision: 84.57792207792207 final fscore: 87.26968174204356  testing metrics: tp = 141, tn = 126, fp = 16, fn = 9 final accuracy: 91.43835616438356 final recall: 94.0 final precision: 89.80891719745223 final fscore: 91.85667752442997</pre>
---	---

### 3-Layer NN using uniform distribution

<pre>training metrics: tp = 393, tn = 432, fp = 158, fn = 185 final accuracy: 70.63356164383562 final recall: 67.99307958477509 final precision: 71.32486388384754 final fscore: 69.6191319751993  testing metrics: tp = 117, tn = 101, fp = 41, fn = 33 final accuracy: 74.65753424657534 final recall: 78.0 final precision: 74.0506329113924 final fscore: 75.97402597402596</pre>	<pre>training metrics: tp = 489, tn = 480, fp = 110, fn = 89 final accuracy: 82.96232876712328 final recall: 84.60207612456747 final precision: 81.63606010016694 final fscore: 83.09260832625318  testing metrics: tp = 130, tn = 115, fp = 27, fn = 20 final accuracy: 83.9041095890411 final recall: 86.66666666666667 final precision: 82.80254777070064 final fscore: 84.69055374592834</pre>
---	--

### 2-Layer NN using normal distribution


<pre>training metrics: tp = 479, tn = 460, fp = 130, fn = 99 final accuracy: 80.39383561643835 final recall: 82.8719723183391 final precision: 78.65353037766832 final fscore: 80.70766638584666  testing metrics: tp = 131, tn = 112, fp = 30, fn = 19 final accuracy: 83.21917808219177 final recall: 87.33333333333333 final precision: 81.36645962732919 final fscore: 84.2443729903537</pre>	<pre>training metrics: tp = 468, tn = 470, fp = 120, fn = 110 final accuracy: 80.3082191780822 final recall: 80.96885813148789 final precision: 79.59183673469387 final fscore: 80.27444253859348  testing metrics: tp = 128, tn = 113, fp = 29, fn = 22 final accuracy: 82.53424657534246 final recall: 85.33333333333334 final precision: 81.52866242038218 final fscore: 83.38762214983714</pre>
---	---

### 3-Layer NN using normal distribution

<pre>training metrics: tp = 519, tn = 492, fp = 98, fn = 59 final accuracy: 86.5582191780822 final recall: 89.79238754325259 final precision: 84.11669367909238 final fscore: 86.86192468619247  testing metrics: tp = 141, tn = 126, fp = 16, fn = 9 final accuracy: 91.43835616438356 final recall: 94.0 final precision: 89.80891719745223 final fscore: 91.85667752442997</pre>	<pre>training metrics: tp = 505, tn = 504, fp = 86, fn = 73 final accuracy: 86.38698630136986 final recall: 87.37024221453287 final precision: 85.44839255499154 final fscore: 86.39863130881095  testing metrics: tp = 138, tn = 127, fp = 15, fn = 12 final accuracy: 90.75342465753424 final recall: 92.0 final precision: 90.19607843137256 final fscore: 91.0891089108911</pre>
---	--

### 3-Layer NN using uniform distribution with momentum and decay

## 9. Additional Results:



Loss vs Epochs for Adam, RMSprop and Momentum respectively

<pre> 100%  training metrics: tp = 504, tn = 511, fp = 79, fn = 74 final accuracy: 86.90068493150685 final recall: 87.19723183391004 final precision: 86.44939965694682 final fscore: 86.82170542635659  testing metrics: tp = 139, tn = 130, fp = 12, fn = 11 final accuracy: 92.12328767123287 final recall: 92.66666666666666 final precision: 92.05298013245033 final fscore: 92.35880398671097 </pre>	<pre> 100%  training metrics: tp = 506, tn = 505, fp = 85, fn = 72 final accuracy: 86.5582191780822 final recall: 87.5432525951557 final precision: 85.6175972927242 final fscore: 86.56971770744225  testing metrics: tp = 137, tn = 126, fp = 16, fn = 13 final accuracy: 90.06849315068493 final recall: 91.33333333333333 final precision: 89.54248366013073 final fscore: 90.42904290429043 </pre>	<pre> 100%  training metrics: tp = 501, tn = 512, fp = 78, fn = 77 final accuracy: 86.72945205479452 final recall: 86.67820069204151 final precision: 86.52849740932642 final fscore: 86.60328435609334  testing metrics: tp = 135, tn = 126, fp = 16, fn = 15 final accuracy: 89.38356164383562 final recall: 90.0 final precision: 89.40397350993378 final fscore: 89.70099667774085 </pre>
--	---	---

3-Layer NN with Adam, RMSprop and Momentum respectively for 3750 epochs